

# TensorFlow in a Nutshell — Part One: Basics



August 22, 2016 by [Camron](#)



Camron

1 Comments

**The fast and easy guide to the most popular Deep Learning framework in the world.**

TensorFlow is a framework created by Google for creating Deep Learning models. Deep Learning is a category of machine learning models that use multi-layer neural networks. The idea of deep learning has been around since 1943 when neurophysiologist Warren McCulloch and mathematician Walter Pitts wrote a paper on how neurons might work and they model a simple neural network using electrical circuits.

Many, many developments have occurred since then. These highly accurate mathematical models are extremely computationally expensive. With recent advances in processing power from GPUs and increasing CPU power Deep Learning has been exploding with popularity.

TensorFlow was created with processing power limitations in mind. Open sourced in November 2015, this library can be ran on computers of all kinds including smartphones. It allows for instant creation of trained production models. It is currently the number 1 Deep Learning framework at the time of writing this article.

Top libraries by Github issues opened		
#1:	2908	BVLC/caffe
#2:	2530	fchollet/keras
#3:	2456	tensorflow/tensorflow
#4:	1801	dmlc/mxnet
#5:	1705	Theano/Theano
#6:	1067	deeplearning4j/deeplearning4j
#7:	693	Microsoft/CNTK
#8:	505	mila-udem/blocks
#9:	498	pfnet/chainer
#10:	494	NVIDIA/DIGITS
#11:	394	Lasagne/Lasagne
#12:	342	torch/torch7
#13:	233	NervanaSystems/neon
#14:	206	tflearn/tflearn
#15:	82	IDSIA/brainstorm
#16:	41	karpathy/convnetjs
#17:	39	amznlabs/amazon-dsstne
#18:	27	torchnet/torchnet

Top libraries by Github stars		
#1:	29967	tensorflow/tensorflow
#2:	11914	BVLC/caffe
#3:	7595	fchollet/keras
#4:	5985	Microsoft/CNTK
#5:	5263	karpathy/convnetjs
#6:	5160	torch/torch7
#7:	4740	dmlc/mxnet
#8:	4316	Theano/Theano
#9:	3723	deeplearning4j/deeplearning4j
#10:	3420	tflearn/tflearn
#11:	3162	amznlabs/amazon-dsstne
#12:	2372	Lasagne/Lasagne
#13:	2149	NervanaSystems/neon
#14:	1577	pfnet/chainer
#15:	1371	NVIDIA/DIGITS
#16:	1147	IDSIA/brainstorm
#17:	870	mila-udem/blocks
#18:	787	torchnet/torchnet

Top libraries by Github contributors		
#1:	348	tensorflow/tensorflow
#2:	244	Theano/Theano
#3:	234	fchollet/keras
#4:	202	BVLC/caffe
#5:	169	dmlc/mxnet
#6:	102	torch/torch7
#7:	84	deeplearning4j/deeplearning4j
#8:	75	Microsoft/CNTK
#9:	72	pfnet/chainer
#10:	50	Lasagne/Lasagne
#11:	48	mila-udem/blocks
#12:	42	NervanaSystems/neon
#13:	39	tflearn/tflearn
#14:	28	NVIDIA/DIGITS
#15:	16	amznlabs/amazon-dsstne
#16:	15	IDSIA/brainstorm
#17:	14	karpathy/convnetjs
#18:	10	torchnet/torchnet

Top libraries by Github forks		
#1:	12506	tensorflow/tensorflow
#2:	7194	BVLC/caffe
#3:	2275	fchollet/keras
#4:	1777	dmlc/mxnet
#5:	1540	Theano/Theano
#6:	1484	torch/torch7
#7:	1291	Microsoft/CNTK
#8:	1264	deeplearning4j/deeplearning4j
#9:	1024	karpathy/convnetjs
#10:	662	Lasagne/Lasagne
#11:	482	amznlabs/amazon-dsstne
#12:	450	NervanaSystems/neon
#13:	412	NVIDIA/DIGITS
#14:	377	pfnet/chainer
#15:	336	tflearn/tflearn
#16:	267	mila-udem/blocks
#17:	161	torchnet/torchnet
#18:	108	IDSIA/brainstorm

Created by Francois Chollet @fchollet (twitter)

## Basic Computational Graph

Everything in TensorFlow is based on creating a computational graph. If you've ever used Theano then this section will look familiar. Think of a computational graph as a network of nodes, with each node known as an operation, running some function that can be as simple as addition or subtraction to as complex as some multi variate equation.

An Operation also referred to as op can return zero or more tensors which can be used later on in the graph. Heres a list of operations with their output for example

```
1  import tensorflow as tf
2
3
4  tf.add(1, 2)
5  # 3
6
7  tf.sub(2, 1)
```

```
8      # 1
9
10     tf.mul(2, 2)
11     # 4
12
13     tf.div(2, 2)
14     # 1
15
16     tf.mod(4, 5)
17     # 4
18
19     tf.pow(3, 2)
20     # 9
21
22     # x < y
23     tf.less(1, 2)
24     # True
25
26     # x <= y
27     tf.less_equal(1, 1)
28     # True
29
30     tf.greater(1, 2)
31     # False
32
33     tf.greater_equal(1, 2)
34     # False
35
36     tf.logical_and(True, False)
37     # False
38
39     tf.logical_or(True, False)
40     # True
41
42     tf.logical_xor(True, False)
43     # True
```

tensorflow\_operations.py hosted with ❤ by [GitHub](#)

[view raw](#)

```
1  import tensorflow as tf
2
3
```

```
4  tf.add(1, 2)
5  # 3
6
7  tf.sub(2, 1)
8  # 1
9
10 tf.mul(2, 2)
11 # 4
12
13 tf.div(2, 2)
14 # 1
15
16 tf.mod(4, 5)
17 # 4
18
19 tf.pow(3, 2)
20 # 9
21
22 # x < y
23 tf.less(1, 2)
24 # True
25
26 # x <= y
27 tf.less_equal(1, 1)
28 # True
29
30 tf.greater(1, 2)
31 # False
32
33 tf.greater_equal(1, 2)
34 # False
35
36 tf.logical_and(True, False)
37 # False
38
39 tf.logical_or(True, False)
40 # True
41
42 tf.logical_xor(True, False)
43 # True
```


Each operation can be handed a constant, array, matrix or n-dimensional matrix. Another word for an n-dimensional matrix is a tensor, a 2-dimensional tensor is equivalent to a m x m matrix.

```
1  import tensorflow as tf
2
3  # create a constant 2X2 matrix
4  tensor_1 = tf.constant([[1., 2.], [3.,4]])
5
6  tensor_2 = tf.constant([[5.,6.],[7.,8.]])
7
8  # create a matrix multiplication operation
9  output_tensor = tf.matmul(tensor_1, tensor_2)
10
11 # have to run the graph using a session
12 sess = tf.Session()
13
14 result = sess.run(output_tensor)
15 print(result)
16
17 sess.close()
```

**tensor\_simple\_operation.py** hosted with ❤ by **GitHub**

[view raw](#)

```
1  import tensorflow as tf
2
3  # create a constant 2X2 matrix
4  tensor_1 = tf.constant([[1., 2.], [3.,4]])
5
6  tensor_2 = tf.constant([[5.,6.],[7.,8.]])
7
8  # create a matrix multiplication operation
9  output_tensor = tf.matmul(tensor_1, tensor_2)
10
11 # have to run the graph using a session
12 sess = tf.Session()
13
14 result = sess.run(output_tensor)
15 print(result)
16
17 sess.close()
```



## Our computational graph

The code above is creating two constant tensors and multiplying them together and outputting our result. This is a trivial example that demonstrates how you can create a graph and run the session. All inputs needed by the op are run automatically. They're typically ran in parallel. This session run actually causes the execution of three operations in the graph, creating the two constants then the matrix multiplication.

## Graph

The constants and operation that we created above was automagically added to the graph in TensorFlow. The graph default is instantiated when the library is imported. Creating a Graph object instead of using the default graph is useful when creating multiple models in one file that do not depend on each other.

```
new_graph = tf.Graph()
```

```
with new_graph.as_default():  
    new_g_const = tf.constant([1., 2.])
```

any variables or operations used outside of the `with new_graph.as_default()` will be added to the default graph that is created when the library is loaded. You can even get a handle to the default graph with

```
default_g = tf.get_default_graph()
```

for most cases it's best to stick with the default graph.

## Session

There are two kinds of Session objects in TensorFlow:

### **tf.Session()**

This encapsulates the environment that operations and tensors are executed and evaluated. Sessions can have their own variables, queues and readers that are allocated. So it's important to use the `close()` method when the session is over. There are 3 arguments for a Session, all of which are optional.

1. `target`—The execution engine to connect to.
2. `graph`—The Graph to be launched.
3. `config`—A ConfigProto protocol buffer with configuration options for the session

To have run one “step” of the TensorFlow computation this function is called and all of the necessary dependencies for the graph to execute are ran.

### **tf.InteractiveSession()**



This is the exact same as `tf.Session()` but is targeted for using IPython and Jupyter Notebooks that allows you to add things and use `Tensor.eval()` and `Operation.run()` instead of having to do `Session.run()` every time you want something to be computed.

```
sess = tf.InteractiveSession()  
a = tf.constant(1)  
b = tf.constant(2)  
c = a + b  
# instead of sess.run(c)  
c.eval()
```

`InteractiveSession` allows so that you don't have to explicitly pass `Session` object.

## Variables

Variables in TensorFlow are managed by the `Session`. They persist between sessions which are useful because `Tensor` and `Operation` objects are immutable. Variables can be created by `tf.Variable()`.

```
tensorflow_var = tf.Variable(1, name="my_variable")
```

most of the time you will want to create these variables as tensors of zeros, ones or random values:

- `tf.zeros()`—creates a matrix full of zeros
- `tf.ones()`—creates a matrix full of ones
- `tf.random_normal()`—a matrix with random uniform values between an interval
- `tf.random_uniform()`—random normally distributed numbers
- `tf.truncated_normal()`—same as random normal but doesn't include any numbers more than 2 standard deviations.

These functions take an initial shape parameter where the dimension of the matrix is defined. For example:



```
# 4x4x4 matrix normally distributed mean 0 std 1  
normal = tf.truncated_normal([4, 4, 4], mean=0.0, stddev=1.0)
```

To have your variable set to one of these matrix helper functions:

```
normal_var = tf.Variable(tf.truncated_normal([4,4,4] , mean=0.0, stddev=1.0))
```

To have these variables initialized you must use TensorFlow's variable initialization function then pass it to the session. This way when multiple sessions are ran the variables are the same.

```
init = tf.initialize_all_variables()  
sess = tf.Session()  
sess.run(init)
```

If you'd like to completely change the value of a variable you can use `Variable.assign()` operation, this must be run in a session update the value.

```
initial_var = tf.Variable(1)
```

```
changed_var = initial_var.assign(initial_var + initial_var)
```

```
init = tf.initialize_all_variables()  
sess = tf.Session()  
sess.run(init)
```

```
sess.run(changed_var)
```

```
# 2
```

```
sess.run(changed_var)
```

```
# 3
```

```
sess.run(changed_var)
```

```
# 4
```

```
# .... and so on
```

Sometimes you would like to add a counter inside your model this is where you can do a `Variable.assign_add()` method which takes a numeric parameter and increments it by the parameter. Similarly there is `Variable.assign_sub()`.

```
counter = tf.Variable(0)
```

```
sess.run(counter.assign_add(1))
```

```
# 1
```

```
sess.run(counter.assign_sub(1))
```

```
# -1
```

## Scope

To control the complexity of models and make them easier to break down into individual pieces TensorFlow has scopes. Scopes are very simple and even help break down your model when using TensorBoard (which will be covered in Part 2). Scopes can even be nested inside of other scopes.

```
with tf.name_scope("Scope1"):  
    with tf.name_scope("Scope_nested"):  
        nested_var = tf.mul(5, 5)
```

Scopes may not seem that powerful right now but used in collaboration with TensorBoard and they're very useful.

## Conclusion

I've demonstrated many of the building blocks that TensorFlow offers. These individual pieces added together can create very complicated models. There is much more that TensorFlow offers, if there are any requests for features in upcoming parts let me know.

It's only fair to share...     

---

*Also published on [Medium](#).*

Tagged [Deep Learning](#), [TensorFlow](#)

## One thought on “TensorFlow in a Nutshell — Part One: Basics”

1. Pingback: [TensorFlow in a Nutshell—Part Two: Hybrid Learning – Camron's Blog](#)

## Leave a Reply

Your email address will not be published. Required fields are marked \*

Comment

Name \*

Email \*

Website

Post Comment

Subscribe

Subscribe to my newsletter for updates and new articles. You won't regret it. A bot may or may not have designed this for success. 10101

Email \*

Subscribe!

Search ...

Recent Posts

- [TensorFlow in a Nutshell—Part Three: All the Models](#)
- [Sales Automation Through a Deep Learning Platform](#)
- [TensorFlow in a Nutshell—Part Two: Hybrid Learning](#)
- [TensorFlow in a Nutshell—Part One: Basics](#)

- [Creating a Search Engine](#)
- 

## Recent Comments

- Camron on [TensorFlow in a Nutshell—Part Two: Hybrid Learning](#)
  - os on [TensorFlow in a Nutshell—Part Two: Hybrid Learning](#)
  - Makis on [TensorFlow in a Nutshell—Part Two: Hybrid Learning](#)
  - Todd Young on [TensorFlow in a Nutshell—Part Two: Hybrid Learning](#)
  - carlo on [TensorFlow in a Nutshell—Part Two: Hybrid Learning](#)
- 

## Categories

- [Uncategorized](#)
- 

Copyright Camron Industries LLC2016



[About Me](#)

Design by Smartcat 

